


Objective

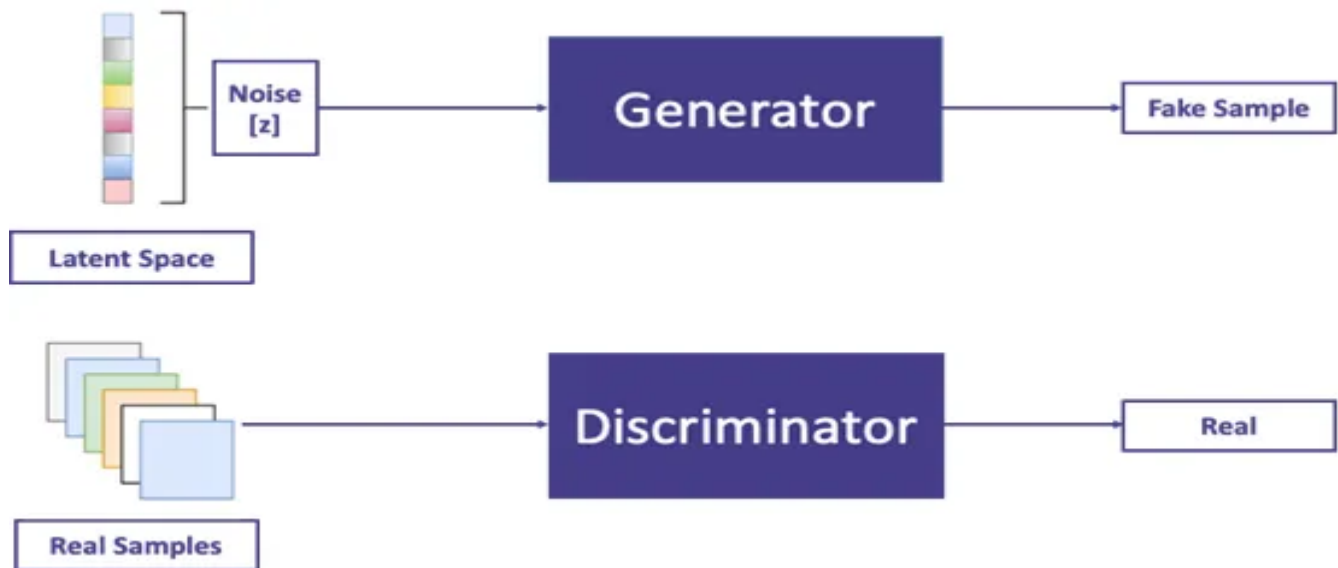
To train a generative adversarial network to generate images using the CIFAR10 dataset.

Basic Architecture for general Overview

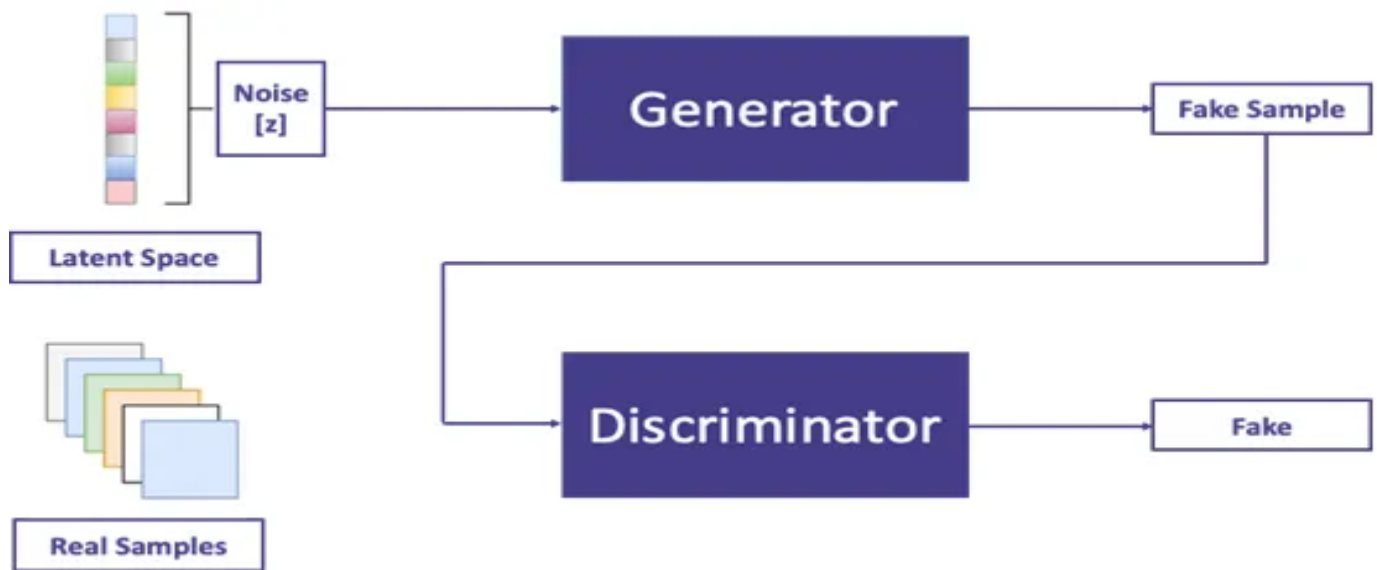
1. During Training



- For fitting real samples during training

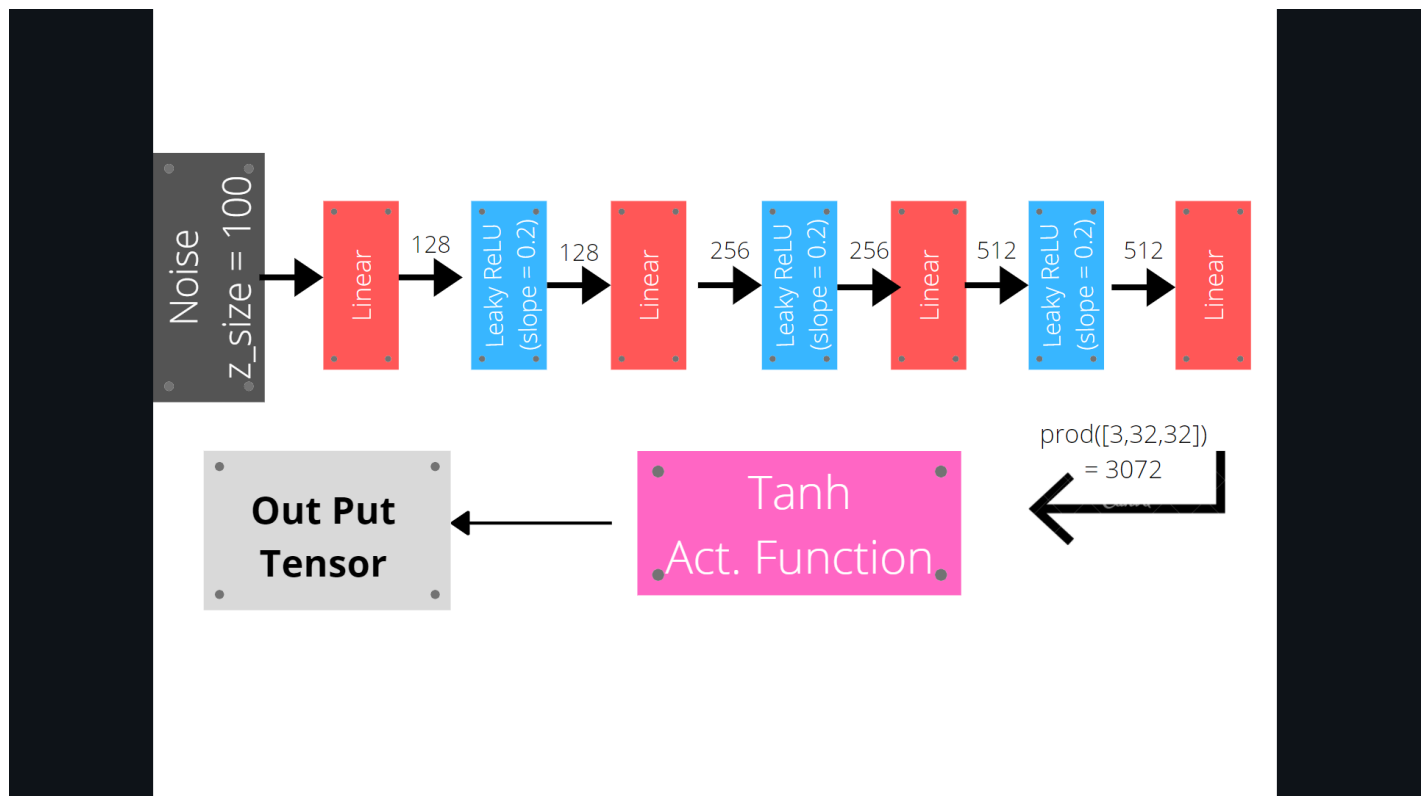


- For fitting fake samples during training

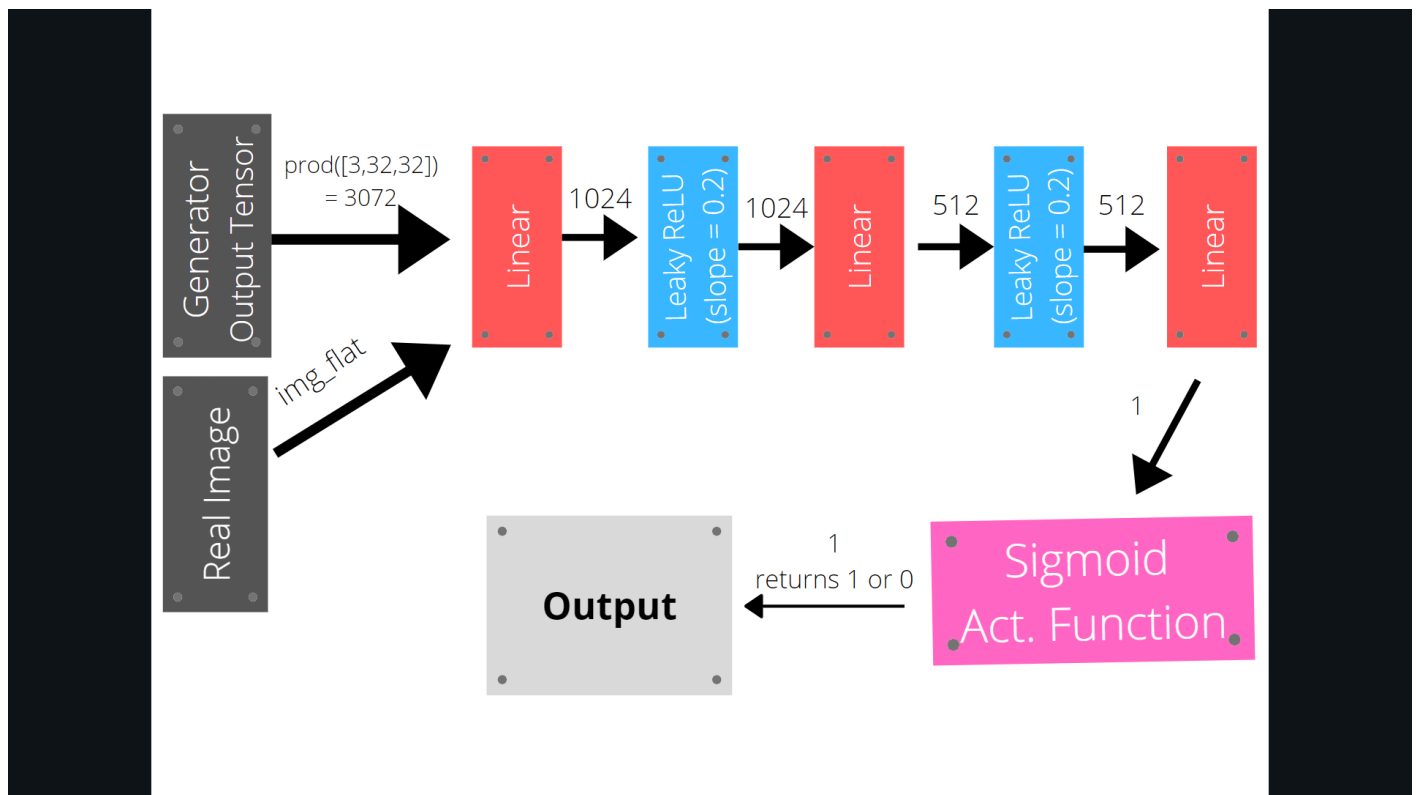


Architecture of Baseline Code

1. Generator Model



2. Discriminator Model



3. Optimisers

- For Generator **SGD** optimiser is used
- For Discriminator Model **RMSprop** is used

4. Hyperparameters

- For updating the parameters **1e-4** is used as the initial learning rate.

5. Batch Size

- In baseline code for training **256** was taken as batch size

Problem Faced

1. No clear image was formed for many epochs.
2. The model was very slow in terms of training i.e taking more epochs to form images.
3. It was taking large epochs to form an image having some indentation of some creature.
4. For each epoch it was taking more time to complete the same.
5. The pictures were not that effective in showing the clear different features of the same.

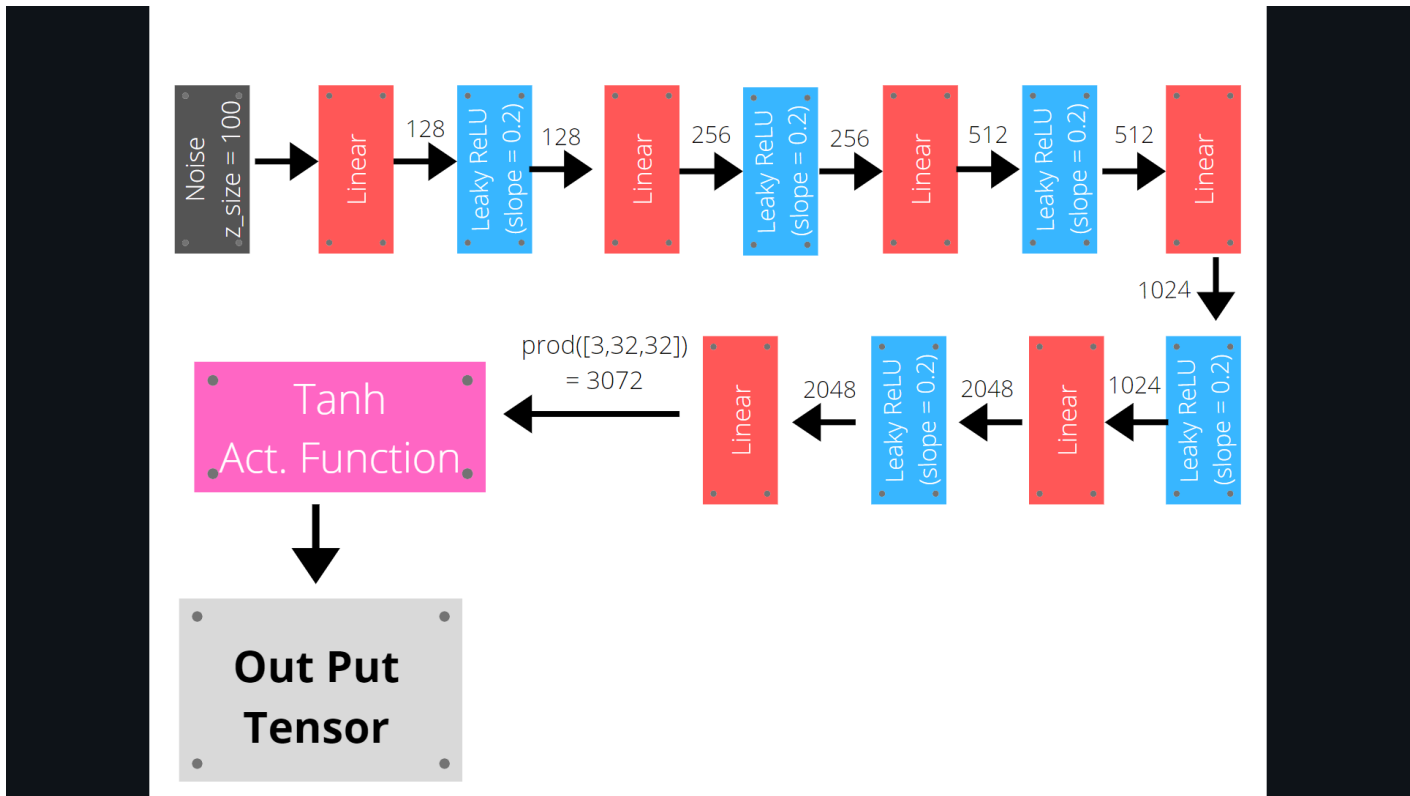
How you overcame them

I have increased the learning rate a bit more. For making GPU's speed more I decreased the batch size a bit so that the GPU can load the batch with more efficiency. So I fixed the same to 128. As the pictures came out were not that effective in showing the clear features of the object I increased the number of layers in both models.

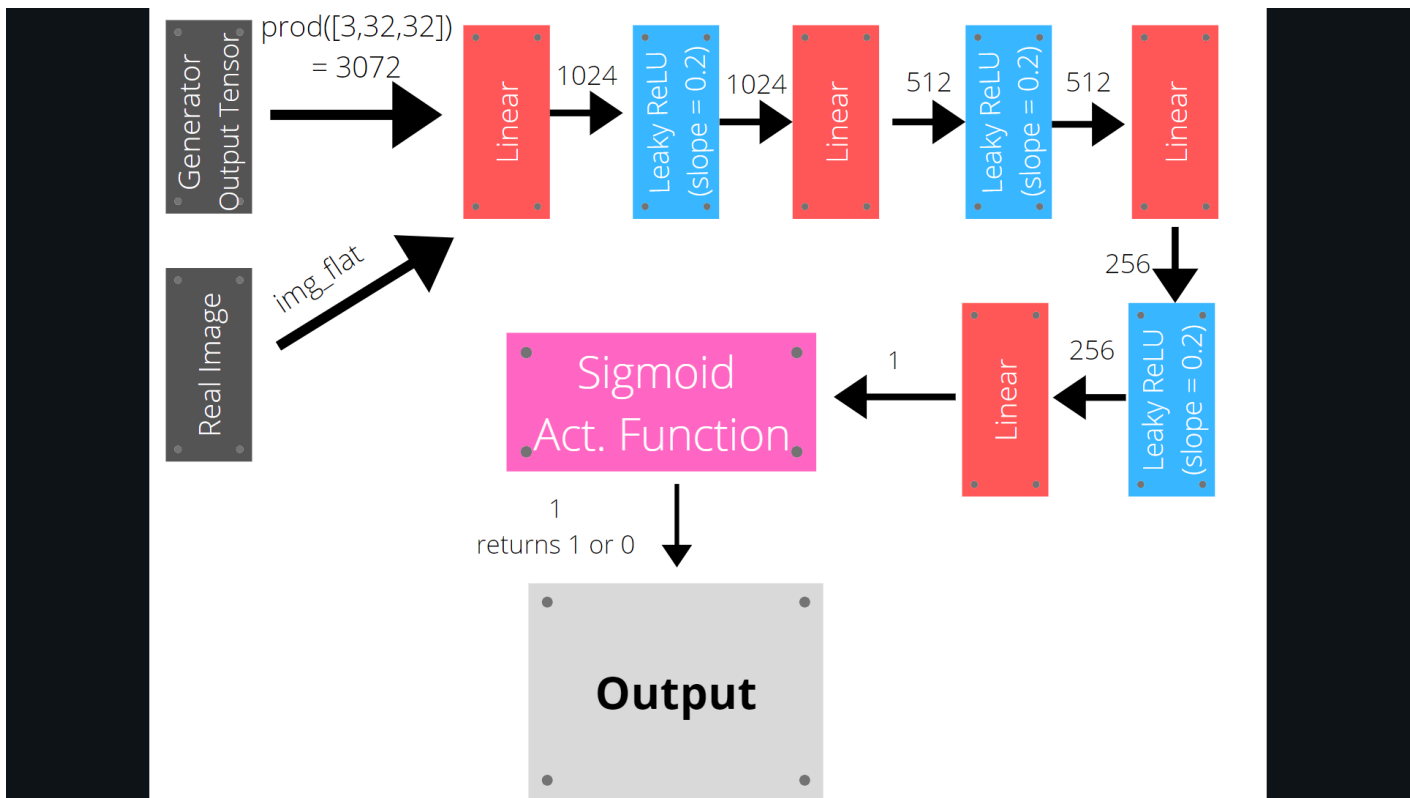
With experimenting with different optimizers finally Adam was the best one for both of my models. Hyperparameters like learning rate had a major impact during training the same, but after shifting from SGD to Adam the training was not that sensitive towards the learning rate. So finally I fixed the learning rate to 0.0002.

Architecture for Final GAN Code

1. Generator Model



2. Discriminator Model



3. Optimisers

- For Generator **Adam** optimiser is used
- For Discriminator Model **Adam** is used

4. Hyperparameters

- For updating the parameters **0.0002** is used as the learning rate.


5. Batch Size

- In final GAN code for training **128** was taken as batch size


Comparison/Improvement

Baseline Model



Epoch = 5



Final Model




Epoch = 10




Baseline Model


Epoch = 20






Epoch = 30



Epoch = 40




Final Model




Baseline Model


Epoch = 80



Final Model




Epoch = 105



...


Epoch = 145




...

Baseline Model

Epoch = 185



Final Model



Here is the loss values I got after 248 epochs

```
+ Code + Text
[248/251][350/391] Loss_D: 1.1595 Loss_G: 1.1262 D(x): 0.3630
[248/251][359/391] Loss_D: 1.1983 Loss_G: 1.6996 D(x): 0.2305
[248/251][360/391] Loss_D: 1.3916 Loss_G: 1.0092 D(x): 0.4173
[248/251][361/391] Loss_D: 1.2786 Loss_G: 1.4237 D(x): 0.2822
[248/251][362/391] Loss_D: 1.1374 Loss_G: 1.2907 D(x): 0.3189
[248/251][363/391] Loss_D: 1.0984 Loss_G: 1.2838 D(x): 0.3305
[248/251][364/391] Loss_D: 1.1298 Loss_G: 1.7003 D(x): 0.2216
[248/251][365/391] Loss_D: 1.1529 Loss_G: 1.2500 D(x): 0.3306
[248/251][366/391] Loss_D: 1.1632 Loss_G: 1.2546 D(x): 0.3314
[248/251][367/391] Loss_D: 1.0837 Loss_G: 1.2998 D(x): 0.3123
[248/251][368/391] Loss_D: 1.1103 Loss_G: 1.5139 D(x): 0.2575
[248/251][369/391] Loss_D: 1.1240 Loss_G: 1.2461 D(x): 0.3321
[248/251][370/391] Loss_D: 1.1530 Loss_G: 1.1101 D(x): 0.3656
[248/251][371/391] Loss_D: 1.2357 Loss_G: 1.3607 D(x): 0.3063
[248/251][372/391] Loss_D: 1.1090 Loss_G: 1.3720 D(x): 0.2915
[248/251][373/391] Loss_D: 1.1986 Loss_G: 1.4375 D(x): 0.2919
[248/251][374/391] Loss_D: 1.1255 Loss_G: 1.2078 D(x): 0.3389
[248/251][375/391] Loss_D: 1.1507 Loss_G: 1.3520 D(x): 0.3030
[248/251][376/391] Loss_D: 1.2299 Loss_G: 1.1573 D(x): 0.3563
[248/251][377/391] Loss_D: 1.1356 Loss_G: 1.3288 D(x): 0.3024
[248/251][378/391] Loss_D: 1.1749 Loss_G: 1.1119 D(x): 0.3819
[248/251][379/391] Loss_D: 1.1685 Loss_G: 1.3552 D(x): 0.3009
[248/251][380/391] Loss_D: 0.9576 Loss_G: 1.4792 D(x): 0.2716
[248/251][381/391] Loss_D: 1.2284 Loss_G: 1.4139 D(x): 0.2843
[248/251][382/391] Loss_D: 1.2254 Loss_G: 1.2403 D(x): 0.3288
[248/251][383/391] Loss_D: 1.1405 Loss_G: 1.2533 D(x): 0.3204
[248/251][384/391] Loss_D: 1.0832 Loss_G: 1.4038 D(x): 0.2887
[248/251][385/391] Loss_D: 1.1744 Loss_G: 1.6208 D(x): 0.2460
[248/251][386/391] Loss_D: 1.1595 Loss_G: 1.0860 D(x): 0.3689
[248/251][387/391] Loss_D: 1.2032 Loss_G: 1.0720 D(x): 0.3786
[248/251][388/391] Loss_D: 1.1739 Loss_G: 1.6041 D(x): 0.2375
[248/251][389/391] Loss_D: 1.2629 Loss_G: 1.0292 D(x): 0.3959
[248/251][390/391] Loss_D: 1.1485 Loss_G: 1.1402 D(x): 0.3530
[248/251][391/391] Loss_D: 1.1427 Loss_G: 1.4515 D(x): 0.2649

The tensorboard extension is already loaded. To reload it, use:
%reload_ext tensorboard
Reusing TensorBoard on port 6006 (pid 1148), started 1:04:48 ago. (Use '!kill 1148' to kill it.)
The tensorboard extension is already loaded. To reload it, use:
%reload_ext tensorboard
```

Insights

After about 200 epochs I have got about to clear images. The GAN algorithm really takes large epochs to be trained. Adjusting hyperparameter & choosing optimizer plays a key role in the efficiency of the model training.

Future Work

The model I made can be made more advanced. There are also GAN algorithms that generate clear images taking much less time in training. I have tried my level best to train my model. There are still some issues which can be checked in my model, So now I am working on that. The GAN models for generating human faces are more complicated, next I'll start learning & trying to make an algorithm for the same.

You can find my code for

- Baseline Model [here](#)
- Final Model [here](#)

Thank you

Omm Prakash Sahoo

Ops#0595

EEE